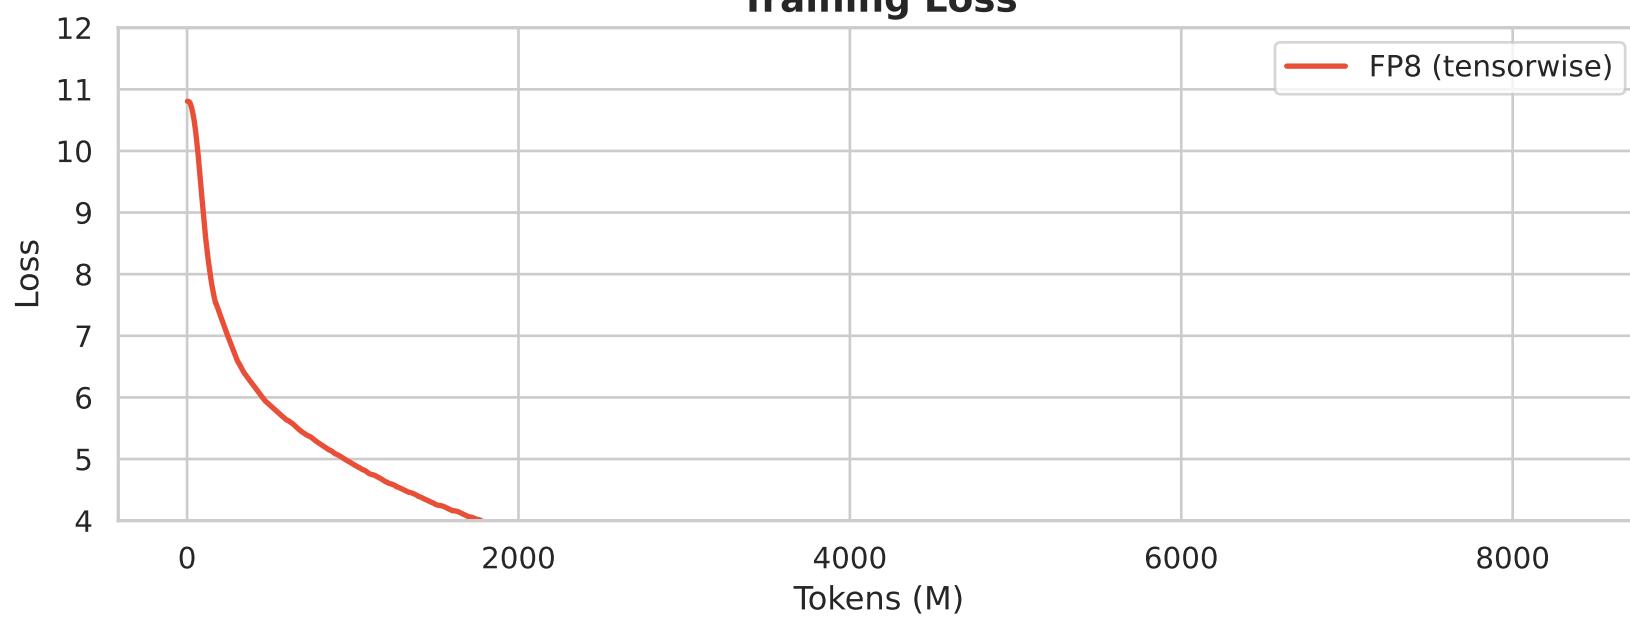


FP8 vs BF16 Training Comparison - 1.5B Parameter Model on NVIDIA GB10

Training Loss

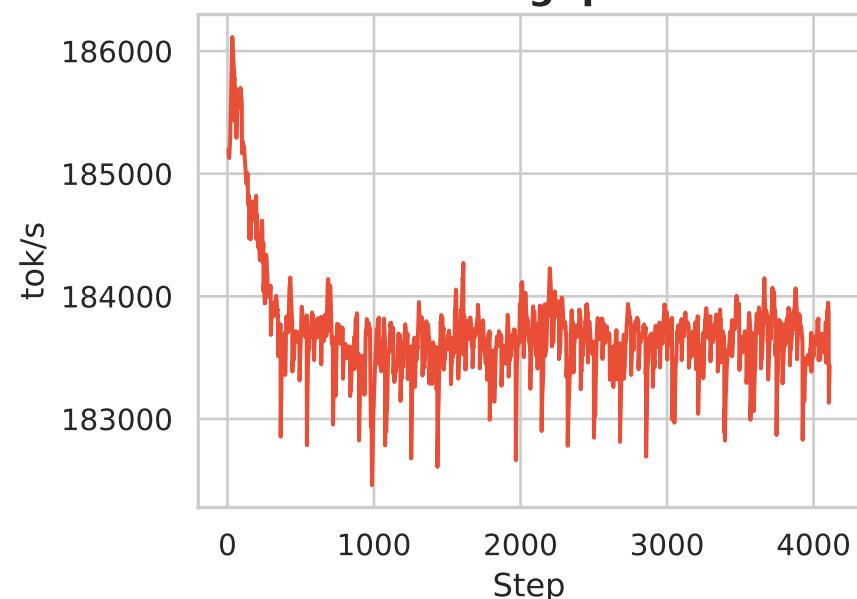


Training Summary

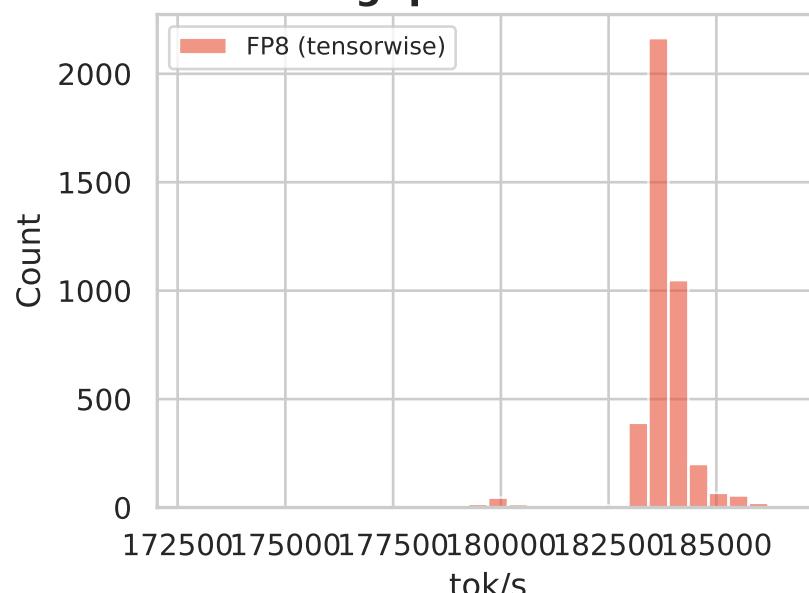
```
=====
```

FP8 (tensorwise)
Final Loss: 2.712
Final PPL: 15.1
Tokens: 8349.9M
Time: 12.69h
Avg tok/s: 183712

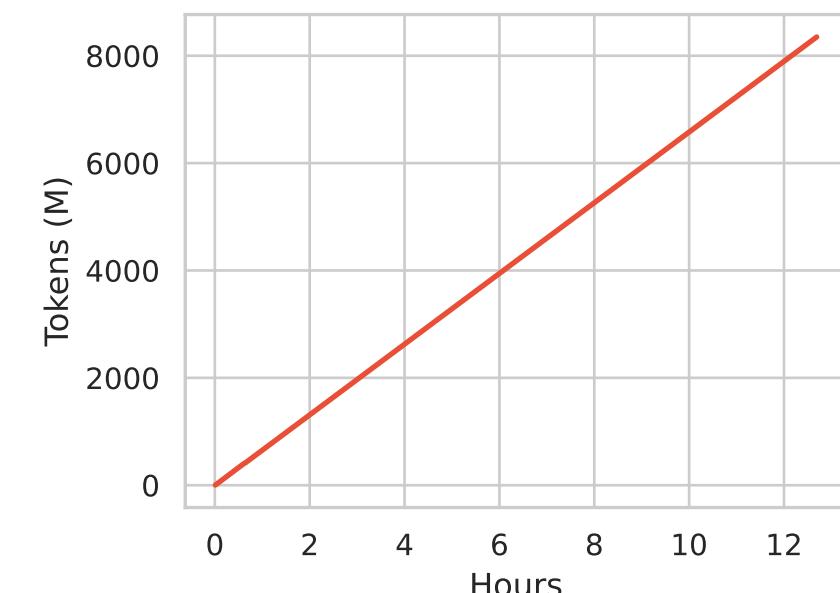
Throughput



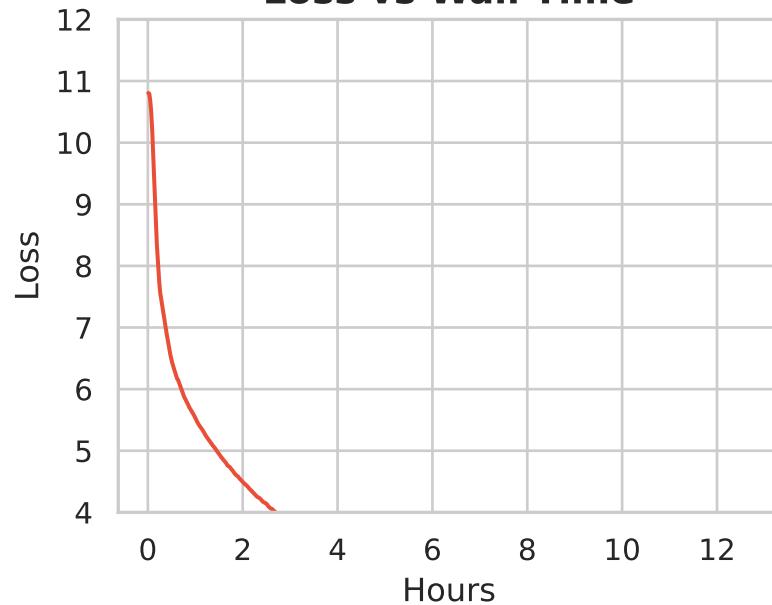
Throughput Distribution



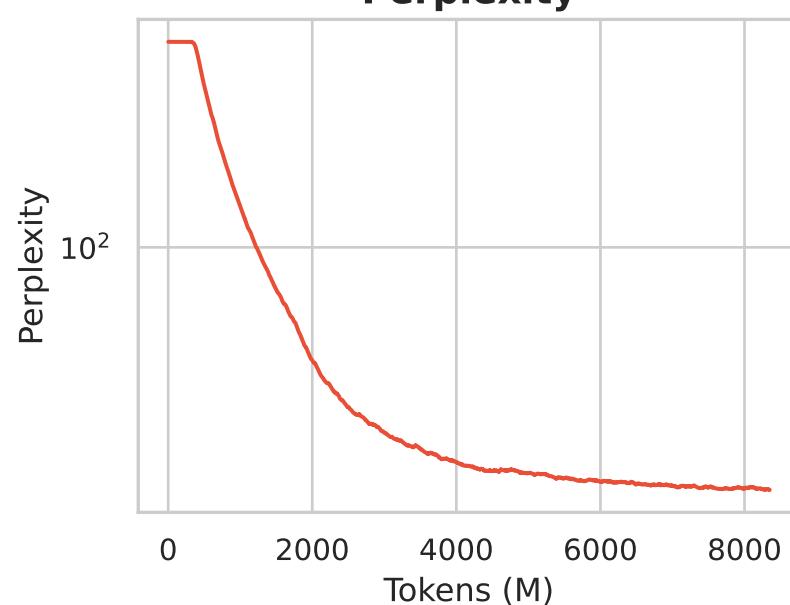
Tokens vs Time



Loss vs Wall Time



Perplexity



Step Time

